

Lecture 3: advanced synchronization primitives

monitor, latch, barrier, thundering herd, semaphore, read-write lock, thread pool, executor,
producer-consumer, fork-join, load balancing

Alexander Filatov
filatovaur@gmail.com

<https://github.com/Svazars/parallel-programming/blob/main/slides/pdf/13.pdf>

In previous episode

We study communication and coordination of different threads in pre-emptive multitasking OS, execution speed is not under our control.

We know the following tools:

- Thread.start, Thread.join, ReentrantLock, Lock, Condition

We are focusing on the following properties:

- Safety, Liveness, Performance

For any new concept we must understand

- Granularity of coordination (race conditions, level of parallelism)
- Blocking API (deadlocks)
- Admission policy (fairness, progress)
- Locking policy (visibility issues, data races)
- Signal send-receive conditions (lost signal, spurious wakeup)

We are focusing on high-level policies rather than particular synchronization mechanisms (until Lecture 11).

Lecture plan

1 Monitor

2 Broadcast

- CountDownLatch
- CyclicBarrier
- Thundering herd problem

3 Group-level concurrency

- Semaphore
- ReadWriteLock

4 Thread pools

5 Workload generation

6 Load balancing

7 Summary

Disclaimer

Concurrent programming is evolving branch of Computer Science.

Disclaimer

Concurrent programming is evolving branch of Computer Science.

There are many designs for key synchronization primitives:

- "Monitor classification" by Buhr et al¹
- "Semaphores in Plan 9" by Mullender and Cox²
- "Futexes are tricky" by Drepper³

¹ <https://dl.acm.org/doi/10.1145/214037.214100>

² <https://swtch.comsemaphore.pdf>

³ <https://www.semanticscholar.org/paper/Futexes-Are-Tricky-Drepper/86253463161f45ea013a501f89c8de36f1a09192>

Disclaimer

Concurrent programming is evolving branch of Computer Science.

There are many designs for key synchronization primitives:

- "Monitor classification" by Buhr et al¹
- "Semaphores in Plan 9" by Mullender and Cox²
- "Futexes are tricky" by Drepper³

We will stick to the Java programming language (built-in monitors) and
java.util.concurrent package.

¹ <https://dl.acm.org/doi/10.1145/214037.214100>

² <https://swtch.comsemaphore.pdf>

³ <https://www.semanticscholar.org/paper/Futexes-Are-Tricky-Drepper/86253463161f45ea013a501f89c8de36f1a09192>

Disclaimer

Concurrent programming is evolving branch of Computer Science.

There are many designs for key synchronization primitives:

- "Monitor classification" by Buhr et al¹
- "Semaphores in Plan 9" by Mullender and Cox²
- "Futexes are tricky" by Drepper³

We will stick to the Java programming language (built-in monitors) and
java.util.concurrent package.

It is enough to illustrate design space, trade-offs, common correctness problems.

¹ <https://dl.acm.org/doi/10.1145/214037.214100>

² <https://swtch.com/semafore.pdf>

³ <https://www.semanticscholar.org/paper/Futexes-Are-Tricky-Drepper/86253463161f45ea013a501f89c8de36f1a09192>

Disclaimer

Concurrent programming is evolving branch of Computer Science.

There are many designs for key synchronization primitives:

- "Monitor classification" by Buhr et al¹
- "Semaphores in Plan 9" by Mullender and Cox²
- "Futexes are tricky" by Drepper³

We will stick to the Java programming language (built-in monitors) and `java.util.concurrent` package.

It is enough to illustrate design space, trade-offs, common correctness problems.

It is **not** enough to start concurrency-oriented low-level development in other languages (e.g. C++ stdlib/boost, Golang with plan9 style sema etc).

¹ <https://dl.acm.org/doi/10.1145/214037.214100>

² <https://swtch.com/sema.pdf>

³ <https://www.semanticscholar.org/paper/Futexes-Are-Tricky-Drepper/86253463161f45ea013a501f89c8de36f1a09192>

Lecture plan

1 Monitor

2 Broadcast

- CountDownLatch
- CyclicBarrier
- Thundering herd problem

3 Group-level concurrency

- Semaphore
- ReadWriteLock

4 Thread pools

5 Workload generation

6 Load balancing

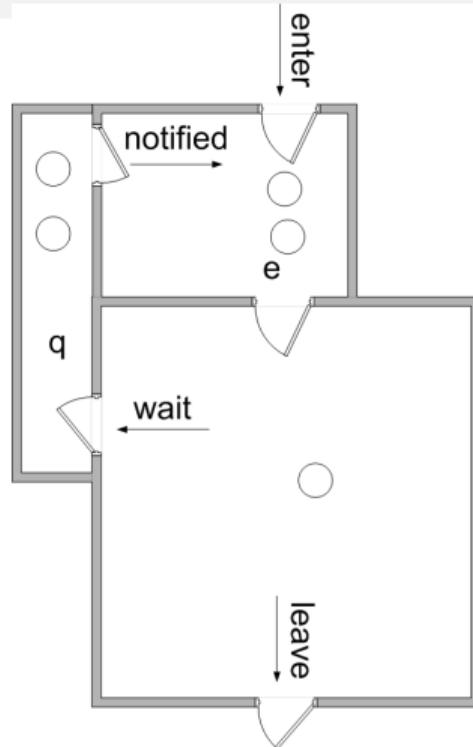
7 Summary

Java monitor

Design

Effectively ReentrantLock + single Condition

- MonitorEnter
- MonitorExit
- wait
- notify
- notifyAll



Java monitor

For every object

Every Java object have associated built-in monitor.

```
Object x = new Object();
synchronized(x) {
    x.wait();
    x.notifyAll();
}
```

Special keyword: synchronized (no .lock/unlock, no try-finally, always structured)

Note: Condition and ReentrantLock are Java objects and they have built-in monitor.

Library-level synchronization and language-level synchronization are **independent**, do not mix notify/signal and wait/await.

Java monitor

For every method

Every Java method could be marked as synchronized.

```
public void foo() { doStuff(); }
public synchronized void bar() { doStuffUnderMutex(); }
```

It synchronizes on this (instance methods) or Class<?> (static methods):

```
public synchronized void fooSignature() { doStuff(); }
public void fooInternal() {
    synchronized(this) {
        doStuff();
    }
}
```

Question time

Question: Is it good or bad to mark methods as thread-safe (synchronized on publicly available signature level?)



Java monitor

For every method

Every Java method could be marked as synchronized.

```
public void foo() { doStuff(); }  
public synchronized void bar() { doStuffunderMutex(); }
```

It synchronizes on this (instance methods) or Class<?> (static methods):

```
public synchronized void fooSignature() { doStuff(); }  
public void fooInternal() {  
    synchronized(this) {  
        doStuff();  
    }  
}
```

Remember that inheritance or synchronized(myAwesomeClassInstance) may break your locking policy and cause a deadlock.

It is common pattern to use private final Object lock and do not expose it as public API.

Java monitor

Single-producer single-consumer bounded queue

```
static Dequeue<Object> buffer = ...  
void producer(Object e) {  
    synchronized(buffer) {  
        while (buffer.size() > N) { buffer.wait(); }  
        buffer.offer(e); buffer.notify();  
    }  
}  
Object consumer() {  
    synchronized(buffer) {  
        while (buffer.isEmpty()) { buffer.wait(); }  
        buffer.notify();  
        return buffer.poll();  
    }  
}
```

Java monitor

Common pitfalls

Question time

Question: Name 3 key properties that you should understand for any concurrent object



Java monitor

Common pitfalls

Correctness(safety):

- Blocking MonitorEnter, wait – deadlock, locking order
- Reentrant
- Visibility/consistency – synchronized looks like "atomic transaction"
- Lost signal
- Predicate invalidation
- Spurious wakeup

Progress(liveness):

- Admission policy for MonitorEnter/MonitorExit is not specified
- Admission policy for wait/notify is not specified

Performance:

- Lock convoy
- Thundering herd

Java monitor

Common pitfalls

Correctness(safety):

- Blocking MonitorEnter, wait – deadlock, locking order
- Reentrant
- Visibility/consistency – synchronized looks like "atomic transaction"
- Lost signal
- Predicate invalidation
- Spurious wakeup

Progress(liveness):

- Admission policy for MonitorEnter/MonitorExit is not specified
- Admission policy for wait/notify is not specified

Performance:

- Lock convoy
- Thundering herd (what is it?)

Java monitor

Specific pitfalls

Unintended encapsulation breach:

- synchronized methods inheritance
- synchronized on object outside of library method

Solution: private guards

Java monitor

Specific pitfalls

Unintended encapsulation breach:

- synchronized methods inheritance
- synchronized on object outside of library method

Solution: private guards

Unintended mix of java.util.concurrent and built-in synchronization:

- wait vs await
- synchronized vs lock/unlock

Solution: static analysis (FindBugs/IntelliJ IDEA)

Java monitor

Rules of thumb

Mutual exclusion

- If you need "just mutex" – consider using encapsulated object + synchronized
 - Save your time with try-finally
 - Other stuff (built-ins have better monitoring and optimizer-friendly) is debatable

Java monitor

Rules of thumb

Mutual exclusion

- If you need "just mutex" – consider using encapsulated object + synchronized
 - Save your time with try-finally
 - Other stuff (built-ins have better monitoring and optimizer-friendly) is debatable
- Use Lock if you need "advanced" properties:
 - Fairness (remember, scheduling is counter-intuitive)
 - Non-structured locking (remember, exceptions are hard)
 - Polling via tryLock (remember, race conditions are tricky)

Java monitor

Rules of thumb

Mutual exclusion

- If you need "just mutex" – consider using encapsulated object + synchronized
 - Save your time with try-finally
 - Other stuff (built-ins have better monitoring and optimizer-friendly) is debatable
- Use Lock if you need "advanced" properties:
 - Fairness (remember, scheduling is counter-intuitive)
 - Non-structured locking (remember, exceptions are hard)
 - Polling via tryLock (remember, race conditions are tricky)

Signalling

Java monitor

Rules of thumb

Mutual exclusion

- If you need "just mutex" – consider using encapsulated object + synchronized
 - Save your time with try-finally
 - Other stuff (built-ins have better monitoring and optimizer-friendly) is debatable
- Use Lock if you need "advanced" properties:
 - Fairness (remember, scheduling is counter-intuitive)
 - Non-structured locking (remember, exceptions are hard)
 - Polling via tryLock (remember, race conditions are tricky)

Signalling. The story is absolutely the same.

Built-in monitors are "friendlier" yet less flexible.

Java monitor

Must-know tool

Monitor is basic universal building block for concurrent primitives and protocols.

Java monitor

Must-know tool

Monitor is basic universal building block for concurrent primitives and protocols.

Homework, mail

*Implement every concurrent primitive presented in this course using built-in Java monitor(s).
This task is **not** graded, but your questions are welcome!*

Java monitor

Must-know tool

Monitor is basic universal building block for concurrent primitives and protocols.

Homework, mail

*Implement every concurrent primitive presented in this course using built-in Java monitor(s).
This task is **not** graded, but your questions are welcome!*

Monitor combines **mutual exclusion** and **signalling** allowing you to make infinite amount of

Java monitor

Must-know tool

Monitor is basic universal building block for concurrent primitives and protocols.

Homework, mail

*Implement every concurrent primitive presented in this course using built-in Java monitor(s).
This task is **not** graded, but your questions are welcome!*

Monitor combines **mutual exclusion** and **signalling** allowing you to make infinite amount of

- concurrent design mistakes

Java monitor

Must-know tool

Monitor is basic universal building block for concurrent primitives and protocols.

Homework, mail

*Implement every concurrent primitive presented in this course using built-in Java monitor(s).
This task is **not** graded, but your questions are welcome!*

Monitor combines **mutual exclusion** and **signalling** allowing you to make infinite amount of

- concurrent design mistakes
- bugs

Java monitor

Must-know tool

Monitor is basic universal building block for concurrent primitives and protocols.

Homework, mail

*Implement every concurrent primitive presented in this course using built-in Java monitor(s).
This task is **not** graded, but your questions are welcome!*

Monitor combines **mutual exclusion** and **signalling** allowing you to make infinite amount of

- concurrent design mistakes
- bugs
- performance regressions

Java monitor

Must-know tool

Monitor is basic universal building block for concurrent primitives and protocols.

Homework, mail

*Implement every concurrent primitive presented in this course using built-in Java monitor(s).
This task is **not** graded, but your questions are welcome!*

Monitor combines **mutual exclusion** and **signalling** allowing you to make infinite amount of

- concurrent design mistakes
- bugs
- performance regressions

Make yourself a favour – learn basics and then progress on advanced stuff.

Lecture plan

1 Monitor

2 Broadcast

- CountDownLatch
- CyclicBarrier
- Thundering herd problem

3 Group-level concurrency

- Semaphore
- ReadWriteLock

4 Thread pools

5 Workload generation

6 Load balancing

7 Summary

Signalling: broadcast

Lock+Condition and Monitor provide notifyAll.

Question time

Question: Does notifyAll or signalAll actually force all waiting threads to wake-up?



Signalling: broadcast

Lock+Condition and Monitor provide notifyAll.

Effectively threads are "released" one-by-one, since they are forced to enter critical section.

Signalling: broadcast

Lock+Condition and Monitor provide notifyAll.

Effectively threads are "released" one-by-one, since they are forced to enter critical section.
What should we do to perform "massive release"?

Signalling: broadcast

Lock+Condition and Monitor provide notifyAll.

Effectively threads are "released" one-by-one, since they are forced to enter critical section.

What should we do to perform "massive release"?

Use specialized tools!

Lecture plan

1 Monitor

2 Broadcast

- CountDownLatch
- CyclicBarrier
- Thundering herd problem

3 Group-level concurrency

- Semaphore
- ReadWriteLock

4 Thread pools

5 Workload generation

6 Load balancing

7 Summary

CountDownLatch



CountDownLatch

CountDownLatch⁴

```
CountDownLatch(int count)
void await()      // awaits counter == 0
void countDown() // counter = max(counter - 1, 0)
long getCount() // return counter
```

Common usage patterns:

- CountDownLatch(1) acts as a "gate":
 - **closed** (counter == 1)
 - **open** (counter == 0)
- CountDownLatch(N) acts as a "all threads ready" point:
 - **not-yet-ready** (counter > 0)
 - **ready** (counter == 0)

⁴

<https://docs.oracle.com/en/java/javase/11/docs/api/java.base/java.util/concurrent/CountDownLatch.html>

CountDownLatch

Example

```
static CountDownLatch ready = new CountDownLatch(N + 1);
static CountDownLatch start = new CountDownLatch(1);
main() {
    for (i = 0; i < N; i++) new Thread(new Worker(i)).start();
    ready.countDown(); ready.await();
    start.countDown();
}
worker() {
    prepare();
    ready.countDown();
    start.await();
}
```

Lecture plan

1 Monitor

2 Broadcast

- CountDownLatch
- CyclicBarrier
- Thundering herd problem

3 Group-level concurrency

- Semaphore
- ReadWriteLock

4 Thread pools

5 Workload generation

6 Load balancing

7 Summary

CyclicBarrier

CountDownLatch is single-use because counter monotonically decreases.

Some tasks require N iterations of some parallel activity before completion.

Question time

Question: Provide examples of parallelizable tasks that consist of N repeating steps.



CyclicBarrier

CountDownLatch is single-use because counter monotonically decreases.

Some tasks require N iterations of some parallel activity before completion.

Creating ArrayList<CountDownLatch> is cumbersome.

CyclicBarrier

CyclicBarrier⁵

```
CyclicBarrier(int parties, Runnable barrierAction)  
int await()  
int getNumberWaiting()  
int getParties()
```

Common usage patterns: mostly-parallel tasks with sequential coordination

- Multiply matrices row-by-row (parallel), handle result (sequential)
- Do numerical simulation for dt (parallel), save intermediate result (sequential)
- Handle batch of requests (parallel), decide next batch size (sequential)

⁵

<https://docs.oracle.com/en/java/javase/11/docs/api/java.base/java.util.concurrent/CyclicBarrier.html>

CyclicBarrier

Example

They have example with matrix multiplication right in the javadoc⁶!

⁶

<https://docs.oracle.com/en/java/javase/11/docs/api/java.base/java/util/concurrent/CyclicBarrier.html>

Lecture plan

1 Monitor

2 Broadcast

- CountDownLatch
- CyclicBarrier
- Thundering herd problem

3 Group-level concurrency

- Semaphore
- ReadWriteLock

4 Thread pools

5 Workload generation

6 Load balancing

7 Summary

Thundering herd problem

```
static Deque<Object> data = ...  
static Lock lock = ...  
static CountDownLatch dataReady = new CountDownLatch(1);  
main(workItems) {  
    lock.lock();  
    try { data.addAll(workItems); } finally { lock.unlock(); }  
    dataReady.countDown();  
}  
worker_i() {  
    dataReady.await();  
    lock.lock();  
    try { Object element = data.poll(); } finally { lock.unlock(); }  
    process(element);  
}
```

Thundering herd problem

```
static Deque<Object> data = ...  
static CountDownLatch dataReady = new CountDownLatch(1);  
main(workItems) {  
    synchronized(data) {  
        data.addAll(workItems);  
    }  
    dataReady.countDown();  
}  
worker_i() {  
    dataReady.await();  
    synchronized(data) {  
        Object element = data.poll();  
    }  
    process(element);  
}
```

Thundering herd problem

Scenario:

- many threads wake-up because of broadcast
- only few (e.g. single one) make progress

Thundering herd problem

Scenario:

- many threads wake-up because of broadcast
- only few (e.g. single one) make progress

Problem: inefficient utilization of scheduling quanta.

Thundering herd problem

Scenario:

- many threads wake-up because of broadcast
- only few (e.g. single one) make progress

Problem: inefficient utilization of scheduling quanta.

Really similar to lock convoy.

Thundering herd problem

Scenario:

- many threads wake-up because of broadcast
- only few (e.g. single one) make progress

Problem: inefficient utilization of scheduling quanta.

Really similar to lock convoy.

Solution:

- Use broadcasting with care
- `monitor.notifyAll` and `condition.signalAll` avoid this by moving signalled thread from one queue to another, waking up single "victim" on `MonitorExit/unlock`⁷.

⁷This is called wait morphing

Thundering herd problem

Scenario:

- many threads wake-up because of broadcast
- only few (e.g. single one) make progress

Problem: inefficient utilization of scheduling quanta.

Really similar to lock convoy.

Solution:

- Use broadcasting with care
- `monitor.notifyAll` and `condition.signalAll` avoid this by moving signalled thread from one queue to another, waking up single "victim" on `MonitorExit/unlock`⁷.

Warm reminder: built-in primitives are good and friendly⁸

⁷This is called wait morphing

⁸High-level overview: "Compact Java Monitors" by Dice, Kogan <https://arxiv.org/pdf/2102.04188.pdf>



Question time

Question: I am confused, could I observe thundering herd problem when using monitors or could not I?



Lecture plan

- 1 Monitor
- 2 Broadcast
 - CountDownLatch
 - CyclicBarrier
 - Thundering herd problem
- 3 Group-level concurrency
 - Semaphore
 - ReadWriteLock
- 4 Thread pools
- 5 Workload generation
- 6 Load balancing
- 7 Summary

Group level concurrency

Mutex divides threads into "owner" and "others".

Condition divides threads into "signalled" and "others".

CountDownLatch/CyclicBarrier help to coordinate start/stop for group of threads.

Could we do more?

Group level concurrency

Mutex divides threads into "owner" and "others".

Condition divides threads into "signalled" and "others".

CountDownLatch/CyclicBarrier help to coordinate start/stop for group of threads.

Could we do more?

- No more than N threads in critical section

Group level concurrency

Mutex divides threads into "owner" and "others".

Condition divides threads into "signalled" and "others".

CountDownLatch/CyclicBarrier help to coordinate start/stop for group of threads.
Could we do more?

- No more than N threads in critical section (semaphore)

Group level concurrency

Mutex divides threads into "owner" and "others".

Condition divides threads into "signalled" and "others".

CountDownLatch/CyclicBarrier help to coordinate start/stop for group of threads.

Could we do more?

- No more than N threads in critical section (semaphore)
- One exclusive writer and N independent readers

Group level concurrency

Mutex divides threads into "owner" and "others".

Condition divides threads into "signalled" and "others".

CountDownLatch/CyclicBarrier help to coordinate start/stop for group of threads.
Could we do more?

- No more than N threads in critical section (semaphore)
- One exclusive writer and N independent readers (read-write lock)

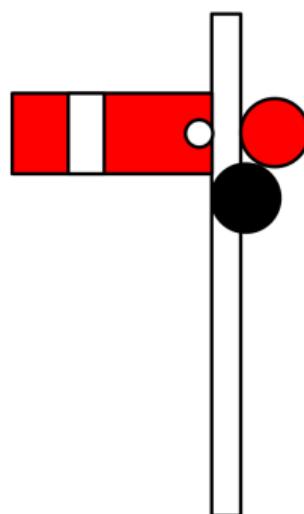
Lecture plan

- 1 Monitor
- 2 Broadcast
 - CountDownLatch
 - CyclicBarrier
 - Thundering herd problem
- 3 Group-level concurrency
 - Semaphore
 - ReadWriteLock
- 4 Thread pools
- 5 Workload generation
- 6 Load balancing
- 7 Summary

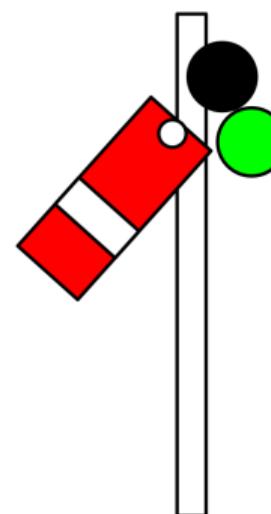
Semaphore

https://en.wikipedia.org/wiki/Railway_semaphore_signal

Stop



Clear



Semaphore

Semaphore⁹

```
Semaphore(int permits)
void acquire() // acquire permit, blocks until available
void release() // release permit
```

Important notes:

- Perfect to
 - limit number of active threads (parallelism control)
 - bound shared resource usage (concurrent access control)
- No ownership – any thread could release permits
- *Binary semaphore* (`permits == 1`) is very similar to `NonReentrantLock`

⁹

<https://docs.oracle.com/en/java/javase/11/docs/api/java.base/java/util/concurrent/Semaphore.html>

Semaphore

Example

```
static Semaphore sema = new Semaphore(N);
static Set<Object> items = ...;
useItem() {
    sema.acquire();
    try {
        synchronized(items) {
            Object item = selectRandom(items);
            items.remove(item);
        }
        use(item);
        synchronized(items) { items.add(item); }
    } finally { sema.release(); }
}
```

Lecture plan

- 1 Monitor
- 2 Broadcast
 - CountDownLatch
 - CyclicBarrier
 - Thundering herd problem
- 3 Group-level concurrency
 - Semaphore
 - **ReadWriteLock**
- 4 Thread pools
- 5 Workload generation
- 6 Load balancing
- 7 Summary

ReadWriteLock

ReadWriteLock¹⁰

```
Lock readLock();  
Lock writeLock();
```

Read-mostly data structures:

- Concurrent readers, Exclusive writer

Design decisions:

¹⁰ <https://docs.oracle.com/en/java/javase/11/docs/api/java.base/java/util/concurrent/locks/ReadWriteLock.html>

Question time

Question: What are the main design decisions for any mutual exclusion primitive?



ReadWriteLock

ReadWriteLock¹¹

Lock `readLock()`

Lock `writeLock()`

Read-mostly data structures:

- Concurrent readers, Exclusive writer

Design decisions:

- Reader preference vs. Writer preference
- Lock upgrade (promotion): `readLock` -> `writeLock`
- Lock downgrade: `writeLock` -> `readLock`
- Reentrancy for writers, reentrancy for readers

Advanced concurrency control: split threads into two types.

¹¹ <https://docs.oracle.com/en/java/javase/11/docs/api/java.base/java/util/concurrent/locks/ReadWriteLock.html>

ReadWriteLock

Example

```
static Map<K, V> kv = ...  
static ReadWriteLock rw = ...  
V get(K k) {  
    rw.readLock().lock();  
    try {  
        return kv.get(k);  
    } finally { rw.readLock().unlock(); }  
}  
void put(K k, V v) {  
    rw.writeLock().lock();  
    try {  
        kv.put(k, v);  
    } finally { rw.writeLock().unlock(); }  
}
```

Lecture plan

1 Monitor

2 Broadcast

- CountDownLatch
- CyclicBarrier
- Thundering herd problem

3 Group-level concurrency

- Semaphore
- ReadWriteLock

4 Thread pools

5 Workload generation

6 Load balancing

7 Summary

Motivation

- Thread is valuable and scarce OS resource.
 - Caching could help.
- Too many Threads may slow-down each other (oversubscription).
 - Soft and hard limits could help.
- Most of the computations do not care which thread is used.
 - Decoupling tasks from carriers could help.
- Threads may have purpose and context (name, priority, task-specific limits).
 - Should be properly reflected in observability API.

ThreadFactory

ThreadFactory¹²

```
Thread newThread(Runnable r)
```

Key concept: abstracting Thread creation, deletion and maintenance.

Smart words: decoupling Thread life cycle management from business logic.

¹²

<https://docs.oracle.com/en/java/javase/11/docs/api/java.base/java/util/concurrent/ThreadFactory.html>

Question time

Question: ThreadFactory that returns "previously used" Thread instance: what are pros and cons?



Executor

Executor¹³

```
void execute(Runnable command)
```

Key concept: abstracting task execution.

Smart words: decoupling business logic (**what** to do) from implementation (**how** and **where** to do) to simplify resource management (**who** will do).

¹³

<https://docs.oracle.com/en/java/javase/11/docs/api/java.base/java/util/concurrent/Executor.html>

Question time

Question: Executor does not return Future. Describe valid Executor implementation that indeed does not require Futures.



ExecutorService

ExecutorService¹⁴

```
void execute(Runnable command) // implements Executor
<T> Future<T> submit(Callable<T> task)
<T> T invokeAny(Collection<Callable<T>> tasks) // returns any successful
void shutdown() // no new tasks
List<Runnable> shutdownNow() // try stop all executing, return all pending
boolean awaitTermination(long timeout, TimeUnit unit)
```

Structured concurrency:

- Asynchronous result (Future)
- Conditional execution (invokeAny)
- Life cycle management (submit, shutdown, awaitTermination)

¹⁴

<https://docs.oracle.com/en/java/javase/11/docs/api/java.base/java.util.concurrent/ExecutorService.html>

Question time

Question: What if somebody submit task with unhandled exception to ExecutorService?
What about other ways to "destroy" thread?



ScheduledExecutorService

ScheduledExecutorService¹⁵

```
// Value-returning one-shot task that becomes enabled after the given delay
<V> ScheduledFuture<V> schedule(Callable<V> callable, long delay,
                                  TimeUnit unit);
// Periodic action that becomes enabled first after the given initial delay
// Executions will start after `initialDelay`, then `initialDelay + period` ...
ScheduledFuture<?> scheduleAtFixedRate(Runnable command, long initialDelay,
                                         long period, TimeUnit unit);
```

Delay and periodic schedule control.

Hint: could be useful for testing, stress measurements and simulating "workload bursts".

¹⁵

<https://docs.oracle.com/en/java/javase/11/docs/api/java.base/java.util.concurrent/ScheduledExecutorService.html>

ThreadPools API: summary

- ThreadFactory – create threads
- Executor – execute tasks
- ExecutorService/ScheduledExecutorService – control task execution (order and resources)

ThreadPools API: summary

- ThreadFactory – create threads
- Executor – execute tasks
- ExecutorService/ScheduledExecutorService – control task execution (order and resources)

Bad news: you do not know which thread will execute your code.

- blocking methods, unexpected lock order, deadlocks, race conditions, data races

ThreadPools API: summary

- ThreadFactory – create threads
- Executor – execute tasks
- ExecutorService/ScheduledExecutorService – control task execution (order and resources)

Bad news: you do not know which thread will execute your code.

- blocking methods, unexpected lock order, deadlocks, race conditions, data races

Good news: you do not know which thread will execute your code.

- under-the-hood synchronization "just works", nothing to implement

ThreadPools API: implementations

Executors¹⁶

```
static ExecutorService newFixedThreadPool(int nThreads, ThreadFactory tFactory)
static ExecutorService newSingleThreadExecutor(ThreadFactory threadFactory);
static ExecutorService newCachedThreadPool(ThreadFactory threadFactory);
```

¹⁶ <https://docs.oracle.com/en/java/javase/11/docs/api/java.base/java/util/concurrent/Executors.html>

¹⁷ <https://docs.oracle.com/en/java/javase/11/docs/api/java.base/java/util/concurrent/ThreadPoolExecutor.html>

ThreadPools API: implementations

Executors¹⁶

```
static ExecutorService newFixedThreadPool(int nThreads, ThreadFactory tFactory)
static ExecutorService newSingleThreadExecutor(ThreadFactory threadFactory);
static ExecutorService newCachedThreadPool(ThreadFactory threadFactory);
```

Remember:

- Fixed-size thread pools may cause resource deadlock
- Single-thread executor not necessarily use the same thread
- Caching can cause memory leaks or unintentionally reuse resources (e.g. native-lib-specific resources)

Full overview and arbitrary customization available¹⁷

¹⁶ <https://docs.oracle.com/en/java/javase/11/docs/api/java.base/java/util/concurrent/Executors.html>

¹⁷ <https://docs.oracle.com/en/java/javase/11/docs/api/java.base/java/util/concurrent/ThreadPoolExecutor.html>

ThreadPools API: now you know it

From now on, you **must** avoid using `new Thread()` in your solutions.

- Use existing ExecutorServices
- Create custom ThreadPoolExecutors
- Encapsulate thread configuration in ThreadFactory

ThreadPools API: now you know it

From now on, you **must** avoid using `new Thread()` in your solutions.

- Use existing `ExecutorServices`
- Create custom `ThreadPoolExecutors`
- Encapsulate thread configuration in `ThreadFactory`

Homework, code

Task 3.1 Rewrite `DiningTable.java`^a using `Executors` to simplify life cycle management (`boolean started`, `boolean shouldStop`).

^a <https://github.com/Svazars/parallel-programming/blob/main/hw/block1/lec2/dining-philosophers/dining-philosophers/src/main/java/org/nsu/syspro/parprog/base/DiningTable.java>

This task is **independent** from dining philosophers problem, you should refactor existing concurrent code, not invent new one.

Lecture plan

1 Monitor

2 Broadcast

- CountDownLatch
- CyclicBarrier
- Thundering herd problem

3 Group-level concurrency

- Semaphore
- ReadWriteLock

4 Thread pools

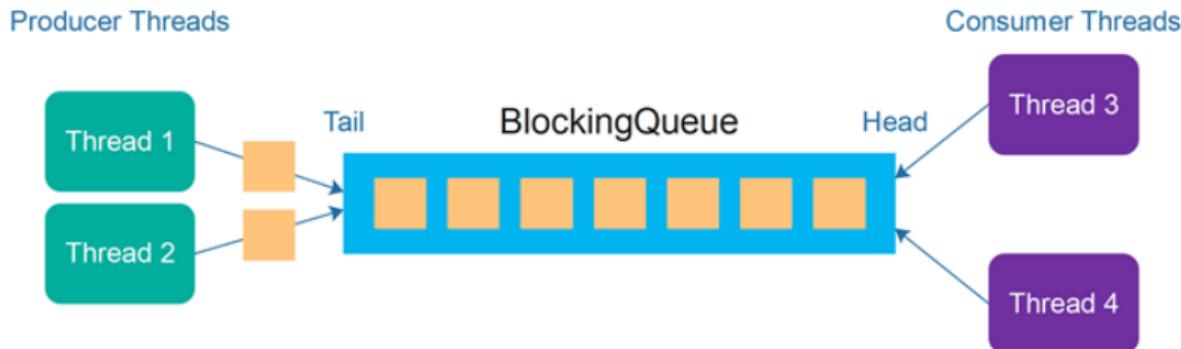
5 Workload generation

6 Load balancing

7 Summary

Producer-consumer

- N threads generate work items
- M threads process work items



BlockingQueue¹⁸

¹⁸

<https://docs.oracle.com/en/java/javase/11/docs/api/java.base/java.util.concurrent/BlockingQueue.html>

Question time

Question: How to ensure that producers and consumers are "balanced"?



Producer-consumer

- X worker threads

Depending on heuristic, every thread either

- generate work item
- process existing work item

Producer-consumer

- X worker threads

Depending on heuristic, every thread either

- generate work item
- process existing work item

Common problems:

- Fixed-size buffer deadlock:
 - buffer capacity is X
 - every thread produces no more than 2 elements per iteration
- Empty-buffer deadlock:
 - buffer is empty
 - every thread heuristically decided to be consumer

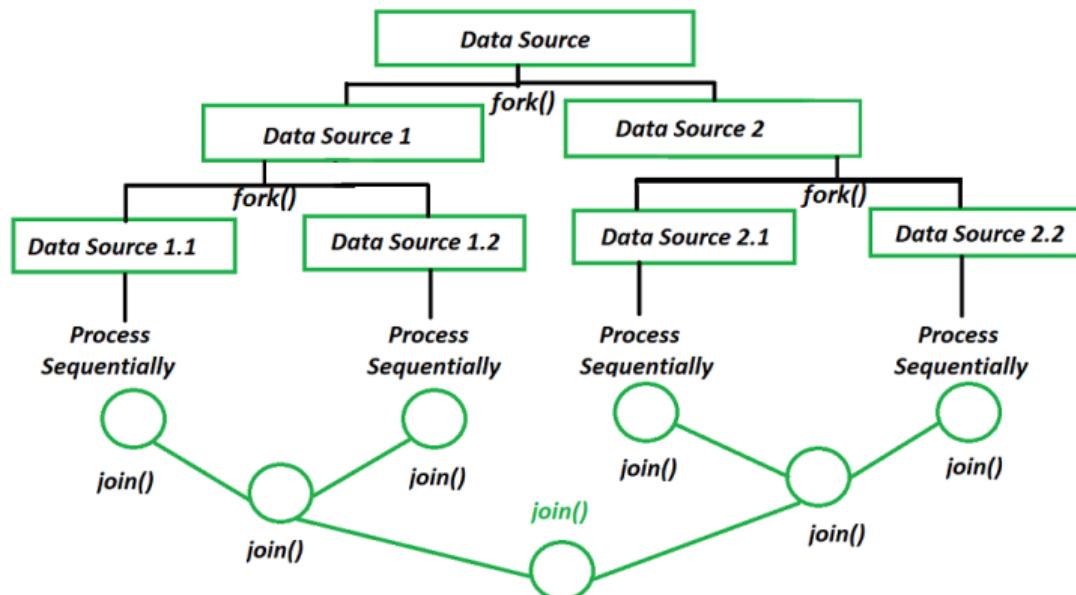
Question time

Question: How to ensure that producers and consumers reach "consistent progress"?



Fork-join

Fork-join model¹⁹



¹⁹https://en.wikipedia.org/wiki/Fork-join_model

Question time

Question: How to decide when to stop fork?



Reactive Streams

<https://openjdk.org/jeps/266>

The main goal of Reactive Streams is to govern the exchange of stream data across an asynchronous boundary—think passing elements on to another thread or thread-pool — while ensuring that the receiving side is not forced to buffer arbitrary amounts of data. In other words, back pressure is an integral part of this model in order to allow the queues which mediate between threads to be bounded.

Flow.Publisher<T>, Flow.Subscriber<T>²⁰

²⁰

<https://docs.oracle.com/en/java/javase/11/docs/api/java.base/java/util/concurrent/Flow.html>

Workload generation: takeaways

- Producer-consumer: "push" based concurrent system
- Fork-join: divide-and-conquer, "helping hand" approach
- Reactive: "propagate backpressure" design

More features – harder to implement, monitor and control.

Lecture plan

1 Monitor

2 Broadcast

- CountDownLatch
- CyclicBarrier
- Thundering herd problem

3 Group-level concurrency

- Semaphore
- ReadWriteLock

4 Thread pools

5 Workload generation

6 Load balancing

7 Summary

Load balancing

ThreadPool-based design uses classic timesharing approach:

- there are X tasks (functions)
- Y "virtual" executors (threads) which are mapped to
- Z "real" executors (cores)

Load balancing

ThreadPool-based design uses classic timesharing approach:

- there are X tasks (functions)
- Y "virtual" executors (threads) which are mapped to
- Z "real" executors (cores)

There are well-known corner cases:

- $X < Y$: insufficient concurrency (coarse-grained problem)
- $X \gg Y$: excessive coordination overheads (too fine-grained approach or too aggressive sub-task creation)
- $Y < Z$: undersubscription (underutilization)
- $Y \gg Z$: oversubscription

Load balancing

ThreadPool-based design uses classic timesharing approach:

- there are X tasks (functions)
- Y "virtual" executors (threads) which are mapped to
- Z "real" executors (cores)

There are well-known corner cases:

- $X < Y$: insufficient concurrency (coarse-grained problem)
- $X \gg Y$: excessive coordination overheads (too fine-grained approach or too aggressive sub-task creation)
- $Y < Z$: undersubscription (underutilization)
- $Y \gg Z$: oversubscription

In any scenario, Y threads have to coordinate task execution:

- Eventual progress: every task will be executed
- Correctness: every task is executed exactly once
- Performance: keep coordination overheads minimal

Load balancing: global queue

Work arbitration

Load balancing design:

- Single shared thread-safe task queue
- Every available Worker gets tasks one-by-one

Counter-examples:

- Many short tasks
- Workers of varying speeds (energy-efficient and usual cores)

Load balancing: global queue + batching

Work arbitration

Load balancing design:

- Single shared thread-safe task queue
- Every available Worker gets tasks in batches (N per request)

Counter-examples:

- Mix of short and long tasks
- Workers of varying speeds (energy-efficient and usual cores)

Load balancing: global queue + batching

Work dealing

Load balancing design:

- Single shared thread-safe task queue
- Every available Worker gets tasks in batches (N per request)
- Overloaded Worker moves fraction of tasks to global queue

Counter-examples:

- Few long tasks
- Overloaded thread spends CPU for auxiliary coordination

Load balancing: work-stealing

Load balancing design:

- Several shared thread-safe task queues (e.g. 1 per Worker)
- Worker get tasks in batches (N per request) from some queue ("steals" from neighbour)

Not that easy to implement and fine-tune²¹. Nontrivial termination protocol.

ForkJoinPool²²

²¹ <https://shipilev.net/#forkjoin>

²² <https://docs.oracle.com/en/java/javase/11/docs/api/java.base/java/util/concurrent/ForkJoinPool.html>

Summary

Mutual exclusion and signalling combined into single universal primitive: monitor
Java language has special support for built-in monitors: synchronized keyword

Bulk thread notification:

- CountDownLatch, CyclicBarrier
- Thundering herd problem

Tools for fine-grained group-level thread control:

- Semaphore, ReadWriteLock

Universal API for thread/task management:

- ThreadFactory, Executor, ExecutorService

Different ways to generate workload concurrently:

- Producer-consumer, Fork-join, Reactive Streams

Designs for efficient concurrent load balancing:

- Work arbitrage, Work dealing, Work stealing

Summary: homework

Homework, mail

*Implement every concurrent primitive presented in this course using built-in Java monitor(s). This task is **not** graded, but your questions are welcome!*

Homework, code

Task 3.1 Rewrite `DiningTable.java`^a using Executors to simplify life cycle management (`boolean started`, `boolean shouldStop`).

^a<https://github.com/Svazars/parallel-programming/blob/main/hw/block1/lec2/dining-philosophers/dining-philosophers/src/main/java/org/nsu/syspro/parprog/base/DiningTable.java>